

# Contents

<b>Preface to the Second Edition</b>	<b>xiii</b>
<b>Preface to the First Edition</b>	<b>xvii</b>
<b>Summary of Notation</b>	<b>xix</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Reinforcement Learning . . . . .	1
1.2 Examples . . . . .	4
1.3 Elements of Reinforcement Learning . . . . .	6
1.4 Limitations and Scope . . . . .	7
1.5 An Extended Example: Tic-Tac-Toe . . . . .	8
1.6 Summary . . . . .	13
1.7 Early History of Reinforcement Learning . . . . .	13
<b>I Tabular Solution Methods</b>	<b>23</b>
<b>2 Multi-armed Bandits</b>	<b>25</b>
2.1 A $k$ -armed Bandit Problem . . . . .	25
2.2 Action-value Methods . . . . .	27
2.3 The 10-armed Testbed . . . . .	28
2.4 Incremental Implementation . . . . .	30
2.5 Tracking a Nonstationary Problem . . . . .	32
2.6 Optimistic Initial Values . . . . .	34
2.7 Upper-Confidence-Bound Action Selection . . . . .	35
2.8 Gradient Bandit Algorithms . . . . .	37
2.9 Associative Search (Contextual Bandits) . . . . .	41
2.10 Summary . . . . .	42

<b>3</b>	<b>Finite Markov Decision Processes</b>	<b>47</b>
3.1	The Agent–Environment Interface . . . . .	47
3.2	Goals and Rewards . . . . .	53
3.3	Returns and Episodes . . . . .	54
3.4	Unified Notation for Episodic and Continuing Tasks . . . . .	57
3.5	Policies and Value Functions . . . . .	58
3.6	Optimal Policies and Optimal Value Functions . . . . .	62
3.7	Optimality and Approximation . . . . .	67
3.8	Summary . . . . .	68
<b>4</b>	<b>Dynamic Programming</b>	<b>73</b>
4.1	Policy Evaluation (Prediction) . . . . .	74
4.2	Policy Improvement . . . . .	76
4.3	Policy Iteration . . . . .	80
4.4	Value Iteration . . . . .	82
4.5	Asynchronous Dynamic Programming . . . . .	85
4.6	Generalized Policy Iteration . . . . .	86
4.7	Efficiency of Dynamic Programming . . . . .	87
4.8	Summary . . . . .	88
<b>5</b>	<b>Monte Carlo Methods</b>	<b>91</b>
5.1	Monte Carlo Prediction . . . . .	92
5.2	Monte Carlo Estimation of Action Values . . . . .	96
5.3	Monte Carlo Control . . . . .	97
5.4	Monte Carlo Control without Exploring Starts . . . . .	100
5.5	Off-policy Prediction via Importance Sampling . . . . .	103
5.6	Incremental Implementation . . . . .	109
5.7	Off-policy Monte Carlo Control . . . . .	110
5.8	*Discounting-aware Importance Sampling . . . . .	112
5.9	*Per-decision Importance Sampling . . . . .	114
5.10	Summary . . . . .	115
<b>6</b>	<b>Temporal-Difference Learning</b>	<b>119</b>
6.1	TD Prediction . . . . .	119
6.2	Advantages of TD Prediction Methods . . . . .	124
6.3	Optimality of TD(0) . . . . .	126
6.4	Sarsa: On-policy TD Control . . . . .	129
6.5	Q-learning: Off-policy TD Control . . . . .	131
6.6	Expected Sarsa . . . . .	133
6.7	Maximization Bias and Double Learning . . . . .	134
6.8	Games, Afterstates, and Other Special Cases . . . . .	136
6.9	Summary . . . . .	138

<b>7</b>	<b><i>n</i>-step Bootstrapping</b>	<b>141</b>
7.1	<i>n</i> -step TD Prediction . . . . .	142
7.2	<i>n</i> -step Sarsa . . . . .	145
7.3	<i>n</i> -step Off-policy Learning . . . . .	148
7.4	*Per-decision Methods with Control Variates . . . . .	150
7.5	Off-policy Learning Without Importance Sampling: The <i>n</i> -step Tree Backup Algorithm . . . . .	152
7.6	*A Unifying Algorithm: <i>n</i> -step $Q(\sigma)$ . . . . .	154
7.7	Summary . . . . .	157
<b>8</b>	<b>Planning and Learning with Tabular Methods</b>	<b>159</b>
8.1	Models and Planning . . . . .	159
8.2	Dyna: Integrated Planning, Acting, and Learning . . . . .	161
8.3	When the Model Is Wrong . . . . .	166
8.4	Prioritized Sweeping . . . . .	168
8.5	Expected vs. Sample Updates . . . . .	172
8.6	Trajectory Sampling . . . . .	174
8.7	Real-time Dynamic Programming . . . . .	177
8.8	Planning at Decision Time . . . . .	180
8.9	Heuristic Search . . . . .	181
8.10	Rollout Algorithms . . . . .	183
8.11	Monte Carlo Tree Search . . . . .	185
8.12	Summary of the Chapter . . . . .	188
8.13	Summary of Part I: Dimensions . . . . .	189
<b>II</b>	<b>Approximate Solution Methods</b>	<b>195</b>
<b>9</b>	<b>On-policy Prediction with Approximation</b>	<b>197</b>
9.1	Value-function Approximation . . . . .	198
9.2	The Prediction Objective ( $\overline{\text{VE}}$ ) . . . . .	199
9.3	Stochastic-gradient and Semi-gradient Methods . . . . .	200
9.4	Linear Methods . . . . .	204
9.5	Feature Construction for Linear Methods . . . . .	210
9.5.1	Polynomials . . . . .	210
9.5.2	Fourier Basis . . . . .	211
9.5.3	Coarse Coding . . . . .	215
9.5.4	Tile Coding . . . . .	217
9.5.5	Radial Basis Functions . . . . .	221
9.6	Selecting Step-Size Parameters Manually . . . . .	222
9.7	Nonlinear Function Approximation: Artificial Neural Networks . . . . .	223
9.8	Least-Squares TD . . . . .	228

9.9	Memory-based Function Approximation . . . . .	230
9.10	Kernel-based Function Approximation . . . . .	232
9.11	Looking Deeper at On-policy Learning: Interest and Emphasis . . . . .	234
9.12	Summary . . . . .	236
<b>10</b>	<b>On-policy Control with Approximation</b>	<b>243</b>
10.1	Episodic Semi-gradient Control . . . . .	243
10.2	Semi-gradient $n$ -step Sarsa . . . . .	247
10.3	Average Reward: A New Problem Setting for Continuing Tasks . . . . .	249
10.4	Deprecating the Discounted Setting . . . . .	253
10.5	Differential Semi-gradient $n$ -step Sarsa . . . . .	255
10.6	Summary . . . . .	256
<b>11</b>	<b>*Off-policy Methods with Approximation</b>	<b>257</b>
11.1	Semi-gradient Methods . . . . .	258
11.2	Examples of Off-policy Divergence . . . . .	260
11.3	The Deadly Triad . . . . .	264
11.4	Linear Value-function Geometry . . . . .	266
11.5	Gradient Descent in the Bellman Error . . . . .	269
11.6	The Bellman Error is Not Learnable . . . . .	274
11.7	Gradient-TD Methods . . . . .	278
11.8	Emphatic-TD Methods . . . . .	281
11.9	Reducing Variance . . . . .	283
11.10	Summary . . . . .	284
<b>12</b>	<b>Eligibility Traces</b>	<b>287</b>
12.1	The $\lambda$ -return . . . . .	288
12.2	TD( $\lambda$ ) . . . . .	292
12.3	$n$ -step Truncated $\lambda$ -return Methods . . . . .	295
12.4	Redoing Updates: Online $\lambda$ -return Algorithm . . . . .	297
12.5	True Online TD( $\lambda$ ) . . . . .	299
12.6	*Dutch Traces in Monte Carlo Learning . . . . .	301
12.7	Sarsa( $\lambda$ ) . . . . .	303
12.8	Variable $\lambda$ and $\gamma$ . . . . .	307
12.9	*Off-policy Traces with Control Variates . . . . .	309
12.10	Watkins's $Q(\lambda)$ to Tree-Backup( $\lambda$ ) . . . . .	312
12.11	Stable Off-policy Methods with Traces . . . . .	314
12.12	Implementation Issues . . . . .	316
12.13	Conclusions . . . . .	317

<b>13 Policy Gradient Methods</b>	<b>321</b>
13.1 Policy Approximation and its Advantages . . . . .	322
13.2 The Policy Gradient Theorem . . . . .	324
13.3 REINFORCE: Monte Carlo Policy Gradient . . . . .	326
13.4 REINFORCE with Baseline . . . . .	329
13.5 Actor–Critic Methods . . . . .	331
13.6 Policy Gradient for Continuing Problems . . . . .	333
13.7 Policy Parameterization for Continuous Actions . . . . .	335
13.8 Summary . . . . .	337
 <b>III Looking Deeper</b>	 <b>339</b>
<b>14 Psychology</b>	<b>341</b>
14.1 Prediction and Control . . . . .	342
14.2 Classical Conditioning . . . . .	343
14.2.1 Blocking and Higher-order Conditioning . . . . .	345
14.2.2 The Rescorla–Wagner Model . . . . .	346
14.2.3 The TD Model . . . . .	349
14.2.4 TD Model Simulations . . . . .	350
14.3 Instrumental Conditioning . . . . .	357
14.4 Delayed Reinforcement . . . . .	361
14.5 Cognitive Maps . . . . .	363
14.6 Habitual and Goal-directed Behavior . . . . .	364
14.7 Summary . . . . .	368
 <b>15 Neuroscience</b>	 <b>377</b>
15.1 Neuroscience Basics . . . . .	378
15.2 Reward Signals, Reinforcement Signals, Values, and Prediction Errors . . . . .	380
15.3 The Reward Prediction Error Hypothesis . . . . .	381
15.4 Dopamine . . . . .	383
15.5 Experimental Support for the Reward Prediction Error Hypothesis . . . . .	387
15.6 TD Error/Dopamine Correspondence . . . . .	390
15.7 Neural Actor–Critic . . . . .	395
15.8 Actor and Critic Learning Rules . . . . .	398
15.9 Hedonistic Neurons . . . . .	402
15.10 Collective Reinforcement Learning . . . . .	404
15.11 Model-based Methods in the Brain . . . . .	407
15.12 Addiction . . . . .	409
15.13 Summary . . . . .	410

<b>16 Applications and Case Studies</b>	<b>421</b>
16.1 TD-Gammon . . . . .	421
16.2 Samuel's Checkers Player . . . . .	426
16.3 Watson's Daily-Double Wagering . . . . .	429
16.4 Optimizing Memory Control . . . . .	432
16.5 Human-level Video Game Play . . . . .	436
16.6 Mastering the Game of Go . . . . .	441
16.6.1 AlphaGo . . . . .	444
16.6.2 AlphaGo Zero . . . . .	447
16.7 Personalized Web Services . . . . .	450
16.8 Thermal Soaring . . . . .	453
<b>17 Frontiers</b>	<b>459</b>
17.1 General Value Functions and Auxiliary Tasks . . . . .	459
17.2 Temporal Abstraction via Options . . . . .	461
17.3 Observations and State . . . . .	464
17.4 Designing Reward Signals . . . . .	469
17.5 Remaining Issues . . . . .	472
17.6 The Future of Artificial Intelligence . . . . .	475
<b>References</b>	<b>481</b>
<b>Index</b>	<b>519</b>